Project 2: Continuous Control

Rory McGrath

February 1, 2019

1 Overview

In this environment, a double-jointed arm acts as the agent and can move to target locations. A reward of +0.1 is provided for each step that the agent's hand is in the goal location. The goal of the agent is to maintain its position at the target location for as many time steps as possible.

The observation space consists of 33 variables corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints. Every entry in the action vector should be a number between -1 and 1.

This task is episodic, in order to solve the environment the agent must achieve an average score of +30 over 100 consecutive episodes.

2 Methods

3 Results

4 Future Work